



Explore-Exploit Graph Traversal for Image Retrieval

Cheng Chang*, Guangwei Yu*, Chundi Liu, Maksims Volkovs.
Layer6 AI

Introduction

- Given a query image, retrieve its most relevant images from an index.



Figure 1: Example query image (left) and four retrieved results (right). Green indicate relevant match while red indicate irrelevant match.

- Motivation:** Relevant images may be visually dissimilar to query, but similar to another image that is similar to query.
- Main Idea:** Traverse the k-nearest neighbor (k-NN) graph to build “trusted” paths between images.




Figure 2: Example retrieval that illustrate our idea. Visually dissimilar image (second in green) is retrieved based on its visually similar neighbors.

k-NN Graph

- Following existing work, we start by building a k-NN graph on the index images.
- Here, vertices correspond to images and edges are weighted by a predefined similarity function.
- We construct the k-NN graph using global descriptor retrieval, and set edge weights to be the dot product between descriptors.




Figure 3: Example k-NN graph with $k=3$. Blue, green, and red denote query, relevant and irrelevant images respectively. Shades of blue denote the neighbors.

Explore-Exploit Graph Traversal (EGT)

- We traverse the k-NN graph starting from query and incrementally build “trusted” paths between vertices by alternating between explore and exploit steps.
- Explore:** neighbors of trusted vertices get added to the explore priority queue ordered by edge weight.
- Exploit:** vertices with edge weights that pass a given threshold get retrieved from the queue and become trusted.
- Edge weight threshold t plays a critical role in our algorithm and controls the degree of exploration.




Figure 4: Example rEGT graph on Google Landmark18. Blue, green and red denote query, relevant, and irrelevant images.




Figure 5: Effect of t on ROxford: at $t=0$ EGT is analogous to breadth-first search (BFS), while $t=\infty$ is analogous to Prim's algorithm.

Algorithm 1: EGT

```

input : k-NN graph  $G_k = (X, A_k, s_k)$ ,
        query  $u$ ,
        number of images to retrieve  $p$ ,
        edge weight threshold  $t$ 
output: list of retrieved images  $Q$ 
1 initialize max-heap  $H$ , list  $V$ , and list  $Q$ 
2 add  $u$  to  $V$ 
3 do
    // Explore step
4   foreach  $v \in V$  do
5     foreach  $x \in NN_k(v), x \notin Q, x \neq u$  do
6       if  $x \in H$  and  $H[x] < s_k(v, x)$  then
7         update weight for  $x$ :  $H[x] \leftarrow s_k(v, x)$ 
8       else if  $x \notin H$  then
9         push  $x$  to  $H$  with weight  $s_k(v, x)$ 
10    end
11  end
12  clear  $V$ 
13  // Exploit step
14  do
15     $v \leftarrow pop(H)$ 
16    add  $v$  to  $V$  and  $Q$ 
17    while ( $peek(H) > t$  or  $|V| = 0$ ) and  $|Q| < p$ 
18    while  $|Q| < p$  and  $|H| > 0$ 
19  return  $Q$ 

```

Edge Re-weighting with Spatial Verification

- k-NN graphs computed from global descriptors (e.g. R-MAC) can be noisy making edge weights unreliable.
- We address this problem by re-weighting edges with spatial verification (rEGT). This significantly reduces topic drift and improves performance.




Figure 6: Performance comparison for rEGT on ROxford Hard.

Efficient Inference

- EGT is conceptually simple and can be efficiently implemented with standard data structures.
- Offline phase only requires k-NN graph construction.
- During online inference, EGT is greedy and independent of the index size, allowing it to scale to large databases.




Figure 7: Performance vs. inference time on ROxford+R1M Hard

Experiments

Table1: Benchmark on ROxford and RParis with and without 1M distractor set, divided into with and without spatial verification (SV).

Method	mAP							
	ROxford		ROxford+R1M		RParis		RParis+R1M	
	Medium	Hard	Medium	Hard	Medium	Hard	Medium	Hard
Without SV								
R-MAC [11]	60.9	32.4	39.3	12.5	78.9	59.4	54.8	28.0
R-MAC+oQE [25]	64.8	36.8	45.7	19.5	82.7	65.7	61.0	35.0
R-MAC+DPS [15]	69.0	44.7	56.6	28.4	89.5	80.0	83.2	70.4
R-MAC+Hybrid-Spectral-Temporal [14]	67.0	44.2	55.6	27.2	89.3	80.2	82.9	69.2
R-MAC+EGT	73.6	56.3	55.8	35.1	90.6	81.2	79.4	63.7
With SV								
HesAff+SIFT+HQE [29]+SV	71.3	49.7	52.0	29.8	70.2	45.1	46.8	21.8
DELFI [20]+HQE+SV	73.4	50.3	60.6	37.9	84.0	69.3	65.2	35.8
HesAffNet+HardNet++ [19]+HQE+SV	75.2	53.3	-	-	73.1	48.9	-	-
R-MAC+DPS+DELFI+ASMK [28]+SV	75.0	48.3	68.7	39.4	90.5	81.2	86.6	74.2
R-MAC+DPS+HesAff+SIFT+ASMK+SV	80.2	54.8	74.9	47.5	92.5	84.0	87.5	76.0
R-MAC+QE+SV+rEGT	83.5	65.8	74.9	54.1	92.8	84.6	87.1	75.6



Figure 8: Top 5 results for diffusion (left) and rEGT (right) from selected queries in ROxford. Blue, green, and red denote query, relevant and irrelevant images respectively.